

```
In [44]: #!/pip install pymysql
#!/pip install --upgrade pip
#!/pip install psycopg2 (for PostgreSQL databases)
#!/pip install kagglehub
#!/pip install psycopg2-binary ipython-sql
import psycopg2 # postgresQL driver for python. To connect.
from sqlalchemy import create_engine, text # text wraps a raw SQL
import os
import pandas as pd
os.environ["KAGGLEHUB_CACHE"] = "/Users/catarina/Desktop/EMDYN"
import kagglehub
import subprocess # very useful module that allows me to more elegantly
from IPython.display import HTML # like the sql, to read html code
```

EMDYN interview

1. Reverts a normal list

```
In [26]: def reverse_list(lst):
        """
        Return a new list that is the reverse of lst.
        """
        return lst[::-1]
```

```
In [27]: # Example:
original = [1, 2, 3, 4, 5]
reversed_list = reverse_list(original)

print("Original:", original)      # Original: [1, 2, 3, 4, 5]
print("Reversed:", reversed_list) # Reversed: [5, 4, 3, 2, 1]

original = ['water', 'fire', 'land', 'air']
reversed_list = reverse_list(original)

print("Original:", original)      # Original: [1, 2, 3, 4, 5]
print("Reversed:", reversed_list) # Reversed: [5, 4, 3, 2, 1]
```

```
Original: [1, 2, 3, 4, 5]
Reversed: [5, 4, 3, 2, 1]
Original: ['water', 'fire', 'land', 'air']
Reversed: ['air', 'land', 'fire', 'water']
```

2. Write a function in Python that reverses a linked list.

A linked list is a fundamental data structure in computer science used to store a sequence of elements—called nodes—where each node holds:

A piece of data (often called its “value” or “payload”). A reference (pointer) to

the next node in the sequence (and in some variants, also to the previous node).

[Node A] → [Node B] → [Node C] → None data data data next next next

Advantages and disadvantages on each list kind

Different. Use array when we need to access entries, when we need to look up by index. When we know the size and it will not increase use array. Linked list: if entries will be deleted. No need to access random indexes.

3. Merge two sorted arrays into a single sorted array

Merge two arrays, A and B. Get the length of each array. Starting from $i=0$ $j=0$. Compare the values of the array for index i and j . Depending on which is larger add to the array C (A+B).

4. Describe how hash tables work, including collision resolution strategies.

Dynamic set of data. They allow for insert/delete/search. Search is the best part because they take in average $O(1)$ (a big O of 1), and a $O(n)$ in the worst case.

NOTE

- $O(1)$ is **constant time**, which means it doesn't take longer as the input size increases. For example, referencing an item in an array takes $O(1)$ time.
- $O(\log n)$ is **logarithmic time**, which means as the input size increases it takes a logarithmically small amount more time. For example, binary searching a sorted list is $O(\log n)$.
- $O(n)$ is **linear time**, which means it takes a constant factor of time proportional to the size of the input size. For example, iterating over every element of an array 5 times is $O(n)$.

Many times confused with a dictionary. A hash table is actually a dictionary using a hash function. Let's look at direct-access tables, similar concept to hash tables. They are actually an array. Constant time operations. If we need to store an infinite number of keys, the inverse is **unbounded and impractical** to store in memory.

In the hash tables the universe is not unbounded, the space is $O(k)$. There is a hash function that maps keys to a location in the table that has data. There

may be two keys with the same data. This **collision** can be dealt with creating a list within the data. The best way is to prevent collisions. A way to do it is with the division. By doing this you create a table with size m and spread your data through the table, with a division. The index of each data stored in the table is obtained by looking at the remainder between the value of the data and the size of the table.

```
In [28]: # Dictionary example
dictionary = {
    'a': 1,
    'b': 9,
    'c': 'C',
    'd': True
}

# insert
dictionary['e'] = False

# delete
del dictionary['a']

# search
print(dictionary['c'])
```

C

5. Write a SQL query to find the second highest salary from a table of employee salaries.

I want to use PostgreSQL.

It is a database management system (RDBMS). It is a software. It allows to store, organize using SQL language. Other alternative could be MySQL and IBM db2.

```
In [29]: # download dataset
path = kagglehub.dataset_download("hummaamqaasim/jobs-in-data")
print("Path to dataset files:", path)
print(os.listdir(path)[0])

file=path+'/'+(os.listdir(path)[0])
print(file)
```

```
Path to dataset files: /Users/catarina/Desktop/EMDYN/datasets/hummaa
mqaasim/jobs-in-data/versions/6
jobs_in_data.csv
/Users/catarina/Desktop/EMDYN/datasets/hummaamqaasim/jobs-in-data/ve
rsions/6/jobs_in_data.csv
```

```
In [30]: # read the file into pandas
df = pd.read_csv(file)
```

```
# show first rows
print("Loaded CSV with shape:", df.shape)
df.head()
```

Loaded CSV with shape: (9355, 12)

Out[30]:

	work_year	job_title	job_category	salary_currency	salary	salary_in_us
--	-----------	-----------	--------------	-----------------	--------	--------------

0	2023	Data DevOps Engineer	Data Engineering	EUR	88000	95000
---	------	----------------------	------------------	-----	-------	-------

1	2023	Data Architect	Data Architecture and Modeling	USD	186000	186000
---	------	----------------	--------------------------------	-----	--------	--------

2	2023	Data Architect	Data Architecture and Modeling	USD	81800	81800
---	------	----------------	--------------------------------	-----	-------	-------

3	2023	Data Scientist	Data Science and Research	USD	212000	212000
---	------	----------------	---------------------------	-----	--------	--------

4	2023	Data Scientist	Data Science and Research	USD	93300	93300
---	------	----------------	---------------------------	-----	-------	-------

In [31]: df.describe(include='all')

Out[31]:

	work_year	job_title	job_category	salary_currency	salary
--	-----------	-----------	--------------	-----------------	--------

count	9355.000000	9355	9355	9355	9355.000000
-------	-------------	------	------	------	-------------

unique	NaN	125	10	11	NaN
--------	-----	-----	----	----	-----

top	NaN	Data Engineer	Data Science and Research	USD	NaN
-----	-----	---------------	---------------------------	-----	-----

freq	NaN	2195	3014	8591	NaN
------	-----	------	------	------	-----

mean	2022.760449	NaN	NaN	NaN	149927.98129
------	-------------	-----	-----	-----	--------------

std	0.519470	NaN	NaN	NaN	63608.83538
-----	----------	-----	-----	-----	-------------

min	2020.000000	NaN	NaN	NaN	14000.00000
-----	-------------	-----	-----	-----	-------------

25%	2023.000000	NaN	NaN	NaN	105200.00000
-----	-------------	-----	-----	-----	--------------

50%	2023.000000	NaN	NaN	NaN	143860.00000
-----	-------------	-----	-----	-----	--------------

75%	2023.000000	NaN	NaN	NaN	187000.00000
-----	-------------	-----	-----	-----	--------------

max	2023.000000	NaN	NaN	NaN	450000.00000
-----	-------------	-----	-----	-----	--------------

In [32]: !brew services start postgresql
!pg_isready
%load_ext sql

Service `postgresql@14` already started, use `brew services restart postgresql@14` to restart.
/tmp:5432 - accepting connections
The sql extension is already loaded. To reload it, use:
%reload_ext sql

```
In [33]: # Function to help with the bash commands, using the subprocess module
def run_shell_command(cmd):
    result = subprocess.run(cmd, capture_output=True, text=True)
    return result

# Create mydata database, or say if it already exists
db_check = run_shell_command([
    "psql", "-U", "postgres", "-tAc",
    "SELECT 1 FROM pg_database WHERE datname='mydata';"
])
if db_check.stdout.strip() == "1":
    print("Database 'mydata' already exists.")
else:
    create_db = run_shell_command(["createdb", "employees"])
    if create_db.returncode == 0:
        print("Database 'mydata' created successfully.")
    else:
        print("Error creating database 'mydata':", create_db.stderr)

# Create my username or say if already exists
user_check = run_shell_command([
    "psql", "-U", "postgres", "-tAc",
    "SELECT 1 FROM pg_roles WHERE rolname='cbranco';"
])
if user_check.stdout.strip() == "1":
    print("User 'cbranco' already exists.")
else:
    create_user = run_shell_command([
        "psql", "-U", "postgres", "-c",
        "CREATE USER cbranco WITH PASSWORD '0000';"
    ])
    if create_user.returncode == 0:
        print("User 'cbranco' created successfully.")
    else:
        print("Error creating user 'cbranco':", create_user.stderr)

# 3) Grant privileges on 'mydata' to 'cbranco'
grant_privs = run_shell_command([
    "psql", "-U", "postgres", "-c",
    "GRANT ALL PRIVILEGES ON DATABASE mydata TO cbranco;"
])
if grant_privs.returncode == 0:
    print("Granted all privileges on 'mydata' to 'cbranco'.")
else:
    print("Error granting privileges:", grant_privs.stderr)
```

Database 'mydata' already exists.
User 'cbranco' already exists.
Granted all privileges on 'mydata' to 'cbranco'.

```
In [34]: # Reuse your Postgres credentials
```

```

DB_USER      = "cbranco"
DB_PASSWORD  = "0000"
DB_HOST      = "localhost"
DB_PORT      = "5432"
DB_NAME      = "mydata"

engine = create_engine(
    f"postgresql+psycopg2://{DB_USER}:{DB_PASSWORD}@{DB_HOST}:{DB_PORT}"
)

#send the df to sql table (employees)
df.to_sql(
    name="employees",
    con=engine,
    if_exists="replace",
    index=False
)

print("✅ CSV has been written into Postgres as table `employees`.")

```

✅ CSV has been written into Postgres as table `employees`.

```

In [35]: # Read back the first 5 rows from Postgres
df2 = pd.read_sql_query("SELECT * FROM employees LIMIT 5;", con=engine)
df2

```

```

Out[35]:

```

	work_year	job_title	job_category	salary_currency	salary	salary_in_usd
0	2023	Data DevOps Engineer	Data Engineering	EUR	88000	95000
1	2023	Data Architect	Data Architecture and Modeling	USD	186000	186000
2	2023	Data Architect	Data Architecture and Modeling	USD	81800	81800
3	2023	Data Scientist	Data Science and Research	USD	212000	212000
4	2023	Data Scientist	Data Science and Research	USD	93300	93300

```

In [36]: sql_snd_dalary = """

SELECT
    salary_in_usd
FROM (
    SELECT
        salary_in_usd,
        DENSE_RANK() OVER (ORDER BY salary_in_usd DESC) AS rnk
    FROM
        employees
) AS ranked_salaries

```

```
WHERE
    rnk = 2;
"""
```

```
In [37]: snd_salary = pd.read_sql_query(sql_snd_dalary, con=engine)
```

```
In [38]: print(snd_salary)
```

```
      salary_in_usd
0             430967
```

```
In [39]: # check which country pays the most
         # to run the sql query, save it to a variable
```

```
sql_highest_salary="""
SELECT
    company_location,
    ROUND(AVG(salary_in_usd)::numeric, 0) AS avg_salary,
    COUNT(*) AS num_jobs
FROM
    employees
GROUP BY
    company_location
ORDER BY
    avg_salary DESC
LIMIT 10;
"""
```

```
In [40]: highest_salary = pd.read_sql_query(sql_highest_salary, con=engine)
```

```
In [41]: print(highest_salary)
```

	company_location	avg_salary	num_jobs
0	Qatar	300000.0	1
1	Puerto Rico	167500.0	4
2	Japan	165500.0	4
3	United States	158159.0	8132
4	Canada	143919.0	226
5	Saudi Arabia	134999.0	2
6	Australia	132283.0	24
7	New Zealand	125000.0	1
8	Ukraine	121333.0	6
9	Bosnia and Herzegovina	120000.0	1

6. Differences between INNER JOIN, LEFT JOIN, RIGHT JOIN, and FULL OUTER JOIN.

- **INNER JOIN** returns only the rows for which there is a match in both tables. In other words, it "intersects" Table A and Table B on the join key(s).
- **LEFT JOIN**, returns all rows from the *left* table (A), plus matched rows from the *right* table (B). If there is no match in B, you still get all the A-

rows, but any columns from B will be NULL.

- **RIGHT JOIN**, returns all rows from the *right* table (B), plus matched rows from the *left* table (A). If there is no match in A, you still get all the B-rows, but any columns from A become NULL.
- **FULL OUTER JOIN**, returns all rows from both tables (A and B). If a row in A has a match in B, you combine them. If a row in A has no match in B, you still see the A-row with NULLs for B's columns. If a row in B has no match in A, you see the B-row with NULLs for A's columns. In practice, you see $A \cap B$ (the overlap), plus A-only, plus B-only.

7. Design a database schema for a social media application's "follow" relationships?

```
In [46]: # create a table for the social media application users and followe
# with the "" I can write exactly how I would write in an SQL query
create_users_followers_tables_sql = ""
-- Create users table
CREATE TABLE IF NOT EXISTS users (
  id          BIGSERIAL      PRIMARY KEY,
  username    VARCHAR(50)    NOT NULL UNIQUE,
  email       VARCHAR(255)   NOT NULL UNIQUE,
  created_at  TIMESTAMP      NOT NULL DEFAULT NOW(),
  password_hash VARCHAR(255) NOT NULL,
  CHECK (username <> '')
);

-- Create follows table
CREATE TABLE IF NOT EXISTS follows (
  follower_id BIGINT      NOT NULL,
  followee_id BIGINT      NOT NULL,
  created_at  TIMESTAMP   NOT NULL DEFAULT NOW(),

  PRIMARY KEY (follower_id, followee_id), -- single combinations, p

  FOREIGN KEY (follower_id)
    REFERENCES users (id)
    ON DELETE CASCADE,
  FOREIGN KEY (followee_id) -- if a user is deleted the whole cascade
    REFERENCES users (id)
    ON DELETE CASCADE,

  CHECK (follower_id <> followee_id)
);

-- Indexes to speed up lookups
CREATE INDEX IF NOT EXISTS idx_follows_follower ON follows (follower_id)
CREATE INDEX IF NOT EXISTS idx_follows_followee ON follows (followee_id)
""

# Execute the DDL in PostgreSQL
with engine.begin() as conn:
```



```
conn.execute(text(create_users_followers_tables_sql))

print("✅ Tables `users` and `follows` have been created (or already exist).")
```

```
In [47]: # create some mockup data
insert_users_sql = """
INSERT INTO users (username, email, password_hash)
VALUES
    ('alice', 'alice@example.com', '$2b$12$abcdefghijklmnopqrstuvwxyz01234567890'),
    ('bob', 'bob@example.com', '$2b$12$mnopqrstuvwxyz345678901234567890'),
    ('carol', 'carol@example.com', '$argon2i$v=19$m=65536,t=2,p=1$abcdefghijklmnopqrstu'),
    ('dave', 'dave@example.com', '$argon2i$v=19$m=65536,t=2,p=1$abcdefghijklmnopqrstu'),
    ('erin', 'erin@example.com', '$2b$12$uvwxyz987654321012345678901234567890'),
ON CONFLICT (username) DO NOTHING;
"""

with engine.begin() as conn:
    conn.execute(text(insert_users_sql))

print("✅ Sample rows inserted into `users`.")
```

```
In [50]: # get the users id, so that then I can fake follow events.
with engine.connect() as conn:
    users_df = pd.read_sql("SELECT id, username FROM users ORDER BY id")

print("\nCurrent users (id ↔ username):")
print(users_df.to_string(index=False))

# build a dictionary where the username is the dictionary key and the user id is the value
user_id_map = dict(zip(users_df["username"], users_df["id"]))

# follow-rows based on those IDs
follows_rows = [
    (user_id_map["alice"], user_id_map["bob"]),
    (user_id_map["alice"], user_id_map["carol"]),
    (user_id_map["bob"], user_id_map["alice"]),
    (user_id_map["carol"], user_id_map["dave"]),
    (user_id_map["carol"], user_id_map["erin"]),
    (user_id_map["erin"], user_id_map["alice"]),
]

# insert them with "ON CONFLICT DO NOTHING" to avoid duplicates
insert_follows = text("""
INSERT INTO follows (follower_id, followee_id)
VALUES (:follower_id, :followee_id)
ON CONFLICT (follower_id, followee_id) DO NOTHING;
""")

with engine.begin() as conn:
    for follower_id, followee_id in follows_rows:
        conn.execute(insert_follows, {"follower_id": follower_id, "followee_id": followee_id})
```

```
print("\n✅ Sample rows inserted into `follows`.")
```

Current users (id ↔ username):

id	username
1	alice
2	bob
3	carol
4	dave
5	erin

✅ Sample rows inserted into `follows`.

```
In [52]: # show
with engine.connect() as conn:
    users_full = pd.read_sql("SELECT * FROM users ORDER BY id;", conn)
    follows_full = pd.read_sql("""
        SELECT
            f.follower_id,
            u1.username AS follower_username,
            f.followee_id,
            u2.username AS followee_username,
            f.created_at
        FROM follows AS f
        JOIN users u1 ON u1.id = f.follower_id
        JOIN users u2 ON u2.id = f.followee_id
        ORDER BY f.follower_id, f.followee_id;
    """, conn)

print("\n--- `users` table: ---")
print(users_full.to_string(index=False))

print("\n--- `follows` table (joined with usernames): ---")
print(follows_full.to_string(index=False))
```

--- `users` table: ---

	id	username	email	created_at	password_hash
1	alice	alice@example.com	2025-06-04 12:11:55.737032	\$2b\$12\$a	bcdefghijk01234567890123456789012345678901234567890
2	bob	bob@example.com	2025-06-04 12:11:55.737032	\$2b\$12\$mno	pqrstuv34567890123456789012345678901234567890123456
3	carol	carol@example.com	2025-06-04 12:11:55.737032	\$argon2	i\$v=19\$m=65536,t=2,p=1\$abcd1234\$xyz9876543210abcdef
4	dave	dave@example.com	2025-06-04 12:11:55.737032	\$argon2i\$	v=19\$m=65536,t=2,p=1\$wxyz5678\$lmn543210abcdef987654
5	erin	erin@example.com	2025-06-04 12:11:55.737032	\$2b\$12\$uv	wxyz98765432101234567890123456789012345678901234567

--- `follows` table (joined with usernames): ---

	follower_id	follower_username	followee_id	followee_username	created_at
1	1	alice	2	bob	2025-06-04 12:20:00.757452
2	1	alice	3	carol	2025-06-04 12:20:00.757452
3	2	bob	1	alice	2025-06-04 12:20:00.757452
4	3	carol	4	dave	2025-06-04 12:20:00.757452
5	3	carol	5	erin	2025-06-04 12:20:00.757452
6	5	erin	1	alice	2025-06-04 12:20:00.757452

```
In [54]: # who is 'alice' following
alice_id = user_id_map['alice']
following_df = pd.read_sql(text("""
    SELECT
        u.id AS user_id,
        u.username AS username,
        f.created_at AS followed_at
    FROM follows f
    JOIN users u ON u.id = f.followee_id
    WHERE f.follower_id = :alice_id
    ORDER BY u.username;
"""), engine, params={"alice_id": alice_id})

print("Users Alice is following:")
print(following_df.to_string(index=False))
```

Users Alice is following:

	user_id	username	followed_at
2	2	bob	2025-06-04 12:20:00.757452
3	3	carol	2025-06-04 12:20:00.757452

8. Describe normalization and denormalization. Give an example where denormalization is preferable.

